

Reading Assignment (week-1)



3 Getting started with neural networks 56

3.1 Anatomy of a neural network 58

Layers: the building blocks of deep learning 58 ▪ Models: networks of layers 59 ▪ Loss functions and optimizers: keys to configuring the learning process 60

3.2 Introduction to Keras 61

Keras, TensorFlow, Theano, and CNTK 62 ▪ Developing with Keras: a quick overview 62

3.3 Setting up a deep-learning workstation 65

Jupyter notebooks: the preferred way to run deep-learning experiments 65 ▪ Getting Keras running: two options 66

*Running deep-learning jobs in the cloud: pros and cons 66
What is the best GPU for deep learning? 66*

3.4 Classifying movie reviews: a binary classification example 68

*The IMDB dataset 68 ▪ Preparing the data 69
Building your network 70 ▪ Validating your approach 73
Using a trained network to generate predictions on new data 76 ▪ Further experiments 77 ▪ Wrapping up 77*

3.5 Classifying newswires: a multiclass classification example 78


*The Reuters dataset 78 ▪ Preparing the data 79
Building your network 79 ▪ Validating your approach 80
Generating predictions on new data 83 ▪ A different way to handle the labels and the loss 83 ▪ The importance of having sufficiently large intermediate layers 83 ▪ Further experiments 84 ▪ Wrapping up 84*

3.6 Predicting house prices: a regression example 85

The Boston Housing Price dataset 85 ▪ Preparing the data 86 ▪ Building your network 86 ▪ Validating your approach using K-fold validation 87 ▪ Wrapping up 91

3.7 Chapter summary 92

Reading Assignment (week-2)



4 *Fundamentals of machine learning* 93

- 4.1 Four branches of machine learning 94
 - Supervised learning* 94 ▪ *Unsupervised learning* 94
 - Self-supervised learning* 94 ▪ *Reinforcement learning* 95
- 4.2 Evaluating machine-learning models 97
 - Training, validation, and test sets* 97 ▪ *Things to keep in mind* 100
- 4.3 Data preprocessing, feature engineering, and feature learning 101
 - Data preprocessing for neural networks* 101 ▪ *Feature engineering* 102
- 4.4 Overfitting and underfitting 104
 - Reducing the network's size* 104 ▪ *Adding weight regularization* 107 ▪ *Adding dropout* 109
- 4.5 The universal workflow of machine learning 111
 - Defining the problem and assembling a dataset* 111
 - Choosing a measure of success* 112 ▪ *Deciding on an evaluation protocol* 112 ▪ *Preparing your data* 112
 - Developing a model that does better than a baseline* 113
 - Scaling up: developing a model that overfits* 114
 - Regularizing your model and tuning your hyperparameters* 114
- 4.6 Chapter summary 116